

Turn plan and step policy

At turn start the harness snapshots a ChatTurnPlan. On every model step it runs prepareReportChatStep and returns the only tools the model may call. Characterization tests lock first-step behavior so a “simpler” change cannot silently reopen a 273-page walk.

1. ChatTurnPlan (once per POST)

Built in assembleChatTurnPlan (turn-plan.ts). Downstream code must consume this object instead of re-classifying the same user text.

Field	Meaning	Source
intent	social · read · write	classifyChatUserIntent then optional Flash-Lite
retrievalPolicy	focused · adaptive · comprehensive	classifyRetrievalPolicy
reviewObjective	Coverage key suffix obj: for document review	planCoverageObjective
alreadyDrafted	Write against a filled/partial section → force read_section first	detectAlreadyDraftedSection
sectionIntent	Section inferred from latest user text (not @ tags)	detectSectionIntentFromText

Latest turn only. Comprehensive shape, inventory-section escalation, and all-scope section intent score the *latest* user message. An earlier “draft equipment” must not keep “draft the remaining sections” on the adaptive/comprehensive path.

2. Intent (what kind of turn is this?)




Figure 1. Empty sections and ready attachments are not a draft request. Agent mode means the model *may* write when asked.

3. Retrieval policy (how hard should we look?)

Policy	When	Harness behavior
focused	Skim, no docs, greeting mapped to no_task, narrow locator	May inject buildAutoEvidence ($\leq 1.5s$, fail-soft). No page walk.
adaptive	Default. Complementary grep + outline	Skip auto-evidence so the model must call search_documents. Hide search after a cited hit, a locate/read, or two empty greps.
comprehensive	Empty inventory / open-set draft with distributed evidence (≥ 12 pages or ≥ 6 outline siblings), explicit every-row language, pushback	Expose start/continue/finish_document_review. Drain remaining pages in one continue (8-wide pool, $\sim 60s$ budget).

C2. A populated table with <date> / <identifier> is placeholder_fill (adaptive), not comprehensive. Empty header-only tables are empty inventories and still need a matching finished review.

4. prepareReportChatStep (every step)

step-policy.ts is the single control plane. The seven former route gates still run in this order; tests require activeTools equality with the old chain.




Figure 2. Early returns win. Do not reorder gates without updating step-policy.test.ts characterization.

5. Locked F1 scenarios

harness-scenarios.ts is the cheap merge gate. A change that drops tokens but changes these first-step tools is a fail — do not “check with a live model.”

id	User text	Expected shape
greeting	hi	social · no tools
sentence_rewrite	change the last sentence in Purpose	write on a filled section · alreadyDrafted · force read_section · no document review
placeholder_fill	fill the placeholders in tables 9-11	adaptive placeholder_fill · populated table with <date> · no start_document_review
empty_inventory	draft Calibration from the attachments	empty ELR table · requireInventoryReview · start a walk whose coverage key matches that section
identifier_lookup	Where is SW-LWB-4 listed?	read · identifier search · no page walk

6. Remaining-sections queue

When the engineer asks to fill several empty draftOrder sections, Agent persists chat_sessions.pending_plan. Each turn drafts the current item (two adjacent non-inventory sections may share a turn). After persist the client POSTs Continue the remaining sections. with autoContinue.

- MJ ELR evidence sections stay in progress until the same turn also drafts the assessment (a count in narrative, plus trend / overallGrade / recommendation).
- Cancel or a new typed prompt pauses the queue. Reopening a thread does not auto-chain.
- A matching finish_document_review this turn unlocks edit_table; do not force another start because the table is still empty — unless the finish was a floor-8 skip that never queued the evidence file (e.g. PRQR).

7. Plan vs Agent vs Ask

Mode	Tools
Plan	Allowlist only: read/list/search/outline/page + document-review + ask_user. New tools missing from PLAN_MODE_CHAT_TOOL_NAMES are silently gone.
Ask	Search/extract only. Write tools stripped from the ToolSet.
Agent	May write. Write tools stay registered on a read turn so a hallucinated edit_table remaps to internal unsupported_tool instead of failing the chat. Keep that fallback out of activeTools.

Andrei chat harness · turn plan and step policy · characterization tests beat Langfuse vibes